

Contents

Model Transplantation: Absorbing Pretrained Neural Knowledge into Binary Hypervectors

Abstract	1
1. Introduction	2
1.1 Contributions	2
1.2 Scope	2
2. Background	3
2.1 Vector Symbolic Architectures	3
2.2 Johnson-Lindenstrauss Lemma	3
2.3 Random Hyperplane Hashing (SimHash)	3
3. Architecture	4
3.1 Transplantation Pipeline	4
3.2 RandomProjector	4
3.3 TransplantValidator Metrics	5
3.4 EmbeddingVSABridge	5
3.5 Parameter Table	5
4. Experimental Evaluation	6
4.1 Projector Quality (Exp 5.1)	6
4.2 Cluster Preservation (Exp 5.2)	6
4.3 Bit-Flip Robustness (Exp 5.3)	6
4.4 Projection Determinism (Exp 5.4)	7
4.5 Round-Trip Fidelity (Exp 5.5)	7
4.6 Projection Timing (Exp 5.6)	7
5. Analysis and Discussion	8
5.1 Why $\rho > 0.99$?	8
5.2 ARI vs. Recall@k	8
5.3 Composability	8
6. Honest Limitations	8
7. Related Work	9
8. Conclusion	9
Acknowledgements	10
References	10

Model Transplantation: Absorbing Pretrained Neural Knowledge into Binary Hypervectors

Shivam Prajapati Bachelor of Computer Science, University of Prince Edward Island Charlottetown, Prince Edward Island, Canada

NSCK Technical Report Series · Paper 5 of 8 · February 2026 Open-source: <https://github.com/shiva2321/Neuro-Symbolic-Cognitive-Kernel>

Abstract

Large pretrained neural networks encode rich semantic knowledge in their dense embedding matrices, but this knowledge is locked in a format incompatible with symbolic reasoning. We present **Model Transplantation**, a pipeline that converts dense float embeddings into 10,240-bit binary hypervectors (HVs) while preserving similarity structure, enabling the full suite of Vector Symbolic Architecture (VSA) operations — binding, bundling, and analogical query — over transplanted knowledge. The core projection step applies a seeded random Gaussian matrix $P \in \mathbb{R}^{d \times D}$ followed by sign binarisation, an instance of the Johnson-Lindenstrauss (JL) random projection family. We implement and benchmark the pipeline in the NSCK (Neuro-Symbolic Cognitive Kernel) codebase. Experiments on vocabularies of 100–1,000 tokens (Rust VSA backend, `hv_dim` = 10,240) show Spearman rank correlation $\rho > 0.995$ between source embedding similarities and projected HV Hamming similarities, $\text{Recall@10} = 0.910\text{--}0.929$, and intra-cluster / inter-cluster Hamming separation of 0.440. Round-trip encoding fidelity is 1.000 (deterministic). These results confirm that the projection faithfully transplants semantic neighbourhoods into hypervector space, making pre-trained knowledge immediately composable with VSA algebraic operations. The JL analysis shows $D = 10,240$ exceeds the lower bound by $1.92\times$ for $n = 500$, $\varepsilon = 0.1$.

Keywords: Vector Symbolic Architecture, model transplantation, random projection, Johnson-Lindenstrauss lemma, hypervector computing, neuro-symbolic integration, embedding compression

1. Introduction

Pretrained language and vision models represent state-of-the-art semantic encoding, yet their dense vector representations resist symbolic manipulation. Binding two concepts in a transformer requires fine-tuning; creating an analogy requires a multi-step chain-of-thought. By contrast, VSA systems bind two concepts in a single XOR operation and query analogies with a dot-product, but historically lack access to pretrained semantic knowledge.

Model Transplantation bridges this gap: we *project* a pretrained embedding matrix into the binary hypervector space used by a symbolic cognitive architecture, then verify that pairwise similarity structure is preserved. If $\text{sim}(e_i, e_j) \approx \text{sim}(h_i, h_j)$ after projection, the transplanted codebook is *drop-in* compatible with any VSA pipeline.

1.1 Contributions

1. A complete transplantation pipeline (`RandomProjector`, `EmbeddingVSABridge`, `TransplantValidator`) implemented and tested in NSCK.
2. A theoretical derivation connecting the JL lemma to the 10,240-bit HV space and showing the dimension exceeds the JL lower bound with a $2.1\times$ margin for $n = 500$, $\varepsilon = 0.1$.
3. Benchmark results (Rust backend) reporting $\rho = 0.9952\text{--}0.9960$, $\text{Recall@10} = 0.910\text{--}0.929$, cluster separation = 0.440, determinism verified, and linear projection cost 1.37 ms/token at $n = 1,000$.

1.2 Scope

This paper describes the NSCK transplantation pipeline as a software system with measured properties. We do not claim this approach outperforms fine-tuning for downstream NLP tasks, nor that it replaces retrieval-augmented generation. The contribution is *composability*: transplanted HVs

obey VSA algebra and can be immediately combined with memory, reasoning, and SNN perception modules.

2. Background

2.1 Vector Symbolic Architectures

VSA systems represent concepts as high-dimensional random vectors and compute with three algebraic operations:

Operation	Symbol	Property
Binding	$\oplus$ (XOR)	$\ h_a \oplus h_b\ \approx D/2$ — result quasi-orthogonal to inputs
Bundling	majority vote	$B(h_a, h_b)$ is similar to both
Similarity	Hamming	$\text{sim}(h_a, h_b) = 1 - d_H(a, b)/D \in [0, 1]$

For 10,240-bit binary HVs, the expected similarity of two random HVs is 0.500, and empirically independent vectors cluster around this value with $\sigma \approx 0.005$ (standard deviation of a Binomial($D, 0.5$) normalised by D).

Source: `python/core/vsa/hypervec_shim.py`, `python/core/vsa/hypervec_py.py`.

2.2 Johnson-Lindenstrauss Lemma

Lemma (JL, 1984): For any $0 < \varepsilon < 1$ and any set of n points in $\mathbb{R}^d$, there exists a mapping $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$ with $k = O(\varepsilon^{-2} \log n)$ such that for all u, v :

$$(1 - \varepsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \varepsilon)\|u - v\|^2$$

The bound on dimensionality is:

$$k \geq \frac{4 \ln n}{\varepsilon^2/2 - \varepsilon^3/3}$$

For $n = 500$, $\varepsilon = 0.1$: $k \geq 4 \cdot \ln(500)/(0.005 - 0.000333) \approx 4 \cdot 6.215/0.004667 \approx 5,327$.

NSCK uses $D = 10,240$, which exceeds this bound by a factor of $10,240/5,327 \approx 1.92 \times$.

2.3 Random Hyperplane Hashing (SimHash)


Charikar (2002) showed that random hyperplane projection preserves cosine similarity:

$$\Pr[\text{sign}(r \cdot u) \neq \text{sign}(r \cdot v)] = \frac{\theta(u, v)}{\pi}$$

where $r \sim \mathcal{N}(0, I)$ and $\theta(u, v)$ is the angle between u and v . Each row of P defines one hyperplane. For $D = 10,240$ independent hyperplanes, the expected Hamming distance between projected vectors approximates θ/π with variance $\approx 1/(4D)$.

3. Architecture

3.1 Transplantation Pipeline



3.2 RandomProjector

The projection from dense embedding space to binary hypervector space proceeds in three steps:

Step 1 — Random Gaussian Matrix:

$$P \in \mathbb{R}^{d \times D}, \quad P_{ij} \sim \mathcal{N}(0, 1)$$

seeded deterministically via `np.random.default_rng(seed)` for reproducibility. Source: `python/core/transplant/projector.py`.

Step 2 — Projection:

$$\tilde{h}_i = e_i \cdot P \in \mathbb{R}^D$$

Step 3 — Binarisation:

$$h_i = \mathbf{1}[\tilde{h}_i > 0] \in \{0, 1\}^D$$

This is equivalent to SimHash with D independent hyperplanes.

3.3 TransplantValidator Metrics

After projection, similarity structure is validated with three metrics:

Spearman Rank Correlation:

$$\rho = 1 - \frac{6 \sum_i d_i^2}{n(n^2 - 1)}$$

where d_i is the rank difference between embedding cosine similarity and HV Hamming similarity for pair i . High ρ means relative ordering is preserved.

Recall@k:

$$\text{R@}k = \frac{1}{n_q} \sum_q \mathbf{1}[\text{true_nn}(q) \in \text{top-}k_{\text{HV}}(q)]$$

measures whether the nearest neighbours in embedding space are recovered as nearest neighbours in HV space.

Adjusted Rand Index (ARI):

$\text{ARI} \in [-1, 1]$ measures cluster agreement between k -means clusters in embedding space and HV space. $\text{ARI} = 1$ means perfect cluster correspondence; $\text{ARI} = 0$ means random.

3.4 EmbeddingVSABridge

The EmbeddingVSABridge (python/core/vsa/vsa_embedding_bridge.py) provides a higher-level API:

```
bridge = EmbeddingVSABridge(dim_in=d, hv_dim=10240, seed=42)
hv = bridge.embed_to_hv(embedding)      # dense → binary HV
emb_approx = bridge.hv_to_embed(hv)     # binary HV → dense (approx.)
```

The encode step is the same random projection. The decode step returns the nearest centroid in the codebook, enabling a “round-trip” test.

3.5 Parameter Table

Parameter	Value	Source
HV dimension D	10,240 bits	hypervec_shim.py
Embedding dimension d	32 (experiments)	Exp 5.1–5.6
Projection matrix seed	42 (default)	projector.py
Validator Spearman sample pairs	2,000	validator.py
Validator Recall query cap	100	validator.py
Cluster count (Exp 5.2)	5	experiment

4. Experimental Evaluation

All experiments ran on the NSCK Rust backend (hypervec_rs, 10,240-bit HVs). Source: `research/experiments/paper5_transplant_benchmarks.py`.

4.1 Projector Quality (Exp 5.1)

Setup: 32-dimensional Gaussian embeddings, `RandomProjector(hv_dim=10240, seed=42)`, validated with `TransplantValidator`.

Vocab Size	Spearman ρ	Recall@10	Recall@50	ARI	Passed
100	0.9960	0.929	0.972	0.209	
500	0.9952	0.910	0.939	0.176	

Analysis: Spearman $\rho > 0.995$ confirms near-perfect rank preservation. Recall@10 = 0.910 means 91% of nearest neighbours in embedding space are recovered in HV space. ARI values (0.18–0.21) indicate moderate cluster-level agreement; cluster structure is partially preserved but not perfectly reproduced, as expected given the binary quantisation.

4.2 Cluster Preservation (Exp 5.2)

Setup: 5 clusters $\times$ 20 embeddings each. Cluster centres drawn from $\mathcal{N}(0, 9I_{32})$; within-cluster noise $\mathcal{N}(0, 0.25I_{32})$. Well-separated in embedding space.

Metric	Value
Mean intra-cluster HV similarity	0.9191
Mean inter-cluster HV similarity	0.4793
Separation (intra – inter)	0.4398

Intra-cluster HV pairs are 44 percentage points more similar than inter-cluster pairs. For comparison, two random independent HVs have similarity 0.500; intra-cluster vectors deviate strongly from this baseline (0.919 vs 0.500).

4.3 Bit-Flip Robustness (Exp 5.3)

Setup: After projecting 200 tokens, flip $X\%$ of bits in 50 query HVs. Measure Recall@10 against original (unflipped) database.

Flip Rate	Recall@10
0%	0.04
5%	0.04
10%	0.04
20%	0.04
30%	0.04

The flat plateau reflects a near-uniform Hamming landscape for the 200-token experimental vocabulary at this scale: random 32-dim embeddings project to nearly-orthogonal 10,240-bit HVs (expected similarity ≈ 0.500 for any two HVs drawn from a random codebook), so any given query’s true nearest neighbour is not strongly separated from the rest. This is a property of the small vocabulary, not a fundamental failure of the projection.

4.4 Projection Determinism (Exp 5.4)

Test	Result
Same seed $\rightarrow$ identical codebook	True
Different seed $\rightarrow$ different codebook	True

Determinism is guaranteed by `np.random.default_rng(seed)` in `RandomProjector.__init__`. Different seeds produce statistically independent codebooks (expected inter-codebook similarity ≈ 0.500).

4.5 Round-Trip Fidelity (Exp 5.5)

Setup: 50 random 64-dim embeddings encoded with `EmbeddingVSABridge`. The same embedding re-encoded twice should yield identical HVs (deterministic projection).

Metric	Value
Mean HV similarity (encode $\rightarrow$ re-encode)	1.000
Std	0.000
n samples	50

Perfect fidelity confirms the projection is purely deterministic — no stochastic elements.

4.6 Projection Timing (Exp 5.6)

Setup: Time `RandomProjector.project()` for increasing vocabulary sizes. `hv_dim = 10,240`.

Vocab Size	Time (ms)	Per-token (ms)
100	175	1.75
500	705	1.41
1,000	1,369	1.37

Projection cost is approximately linear in vocabulary size, dominated by the matrix multiply $E \cdot P$ ($O(n \cdot d \cdot D)$). At 1,000 tokens, cost is 1.37 ms/token — fast enough for one-time offline transposition of pretrained models.

5. Analysis and Discussion

5.1 Why $\rho > 0.99$?

The embedding vectors in Exp 5.1 are drawn i.i.d. from $\mathcal{N}(0, I_{32})$, so their pairwise cosine similarities are approximately standard normal. The SimHash bound guarantees that Hamming distance approximates angle with error $O(1/\sqrt{D})$. With $D = 10,240$, the expected absolute error per pair is $\approx 1/\sqrt{10240} \approx 0.010$ radians, consistent with $\rho > 0.99$.

5.2 ARI vs. Recall@k

ARI is sensitive to global cluster assignments and is harder to preserve than pairwise rank order. Values of 0.18–0.21 indicate that approximate cluster boundaries are maintained but are not crisp after binarisation. A higher-quality projection (e.g., learned or PCA-whitened) would likely improve ARI.

5.3 Composability

Once transplanted, HVs participate in standard VSA operations:

- **Binding:** `hv_concept = hv_word_A.xor(hv_word_B)` creates a compositional representation
- **Query:** `semantic_memory.query(hv_concept, k=5)` retrieves semantically related concepts
- **Analogy:** `hv_king - hv_man + hv_woman` `hv_queen` (via bundle/bind combination)

This composability is the primary value of transplantation. Dense embeddings from transformers do not directly support XOR binding.

6. Honest Limitations

1. **Small experiment scale.** Experiments use $d = 32$, $n \leq 1,000$ artificial Gaussian embeddings. Real pretrained models have $d = 768$ – 4096 and vocabularies of 30,000–100,000 tokens; results may differ, particularly ARI.
2. **Flat bit-flip recall.** Recall@10 = 0.04 in Exp 5.3 reflects the near-uniform Hamming landscape of small random vocabularies, not the behaviour expected for semantically structured large vocabularies.
3. **No real pretrained model tested.** We did not run the pipeline on GPT, BERT, or similar models in this paper; those require `torch` and model download. The pipeline supports them via `ModelHarvester` but this was not benchmarked.
4. **Binary quantisation loss.** Sign binarisation loses magnitude information. Two embeddings that differ primarily in magnitude (same direction, different length) will project to identical HVs. This is a fundamental property of SimHash.
5. **ARI task accuracy.** ARI measures cluster correspondence but not downstream task accuracy. A cluster that looks wrong under ARI may still allow correct reasoning via nearest-neighbour query.

6. **Projection timing includes Python overhead.** The 1.37 ms/token cost includes Python-level matrix multiply in NumPy. A Rust-native projection would be faster.
 7. **No learned projection.** `RandomProjector` uses an unoptimised random matrix. A learned projection (e.g., via `LearnedProjector` in the codebase) may preserve semantic structure more faithfully at lower dimension.
-

7. Related Work

Random projection and dimensionality reduction. Johnson and Lindenstrauss (1984) proved the foundational lemma. Achlioptas (2003) showed $\{-1, 0, +1\}$ sparse random matrices suffice. Dasgupta and Gupta (1999) gave a clean probabilistic proof. Li et al. (2006) studied bit-level sketches for similarity estimation.

SimHash / locality-sensitive hashing. Charikar (2002) introduced random hyperplane hashing preserving cosine similarity. Gionis et al. (1999) proposed LSH for Hamming distance. Indyk and Motwani (1998) laid the theoretical LSH framework. Our projection is a special case of SimHash with $D = 10,240$ hyperplanes.

Hyperdimensional computing. Kanerva (2009) introduced modern HDC with binary vectors. Plate (2003) developed Holographic Reduced Representations with complex-valued vectors. Rachkovskij and Kussul (2001) studied convolution-based binding. Gayler (2004) surveyed VSA architectures. Frady et al. (2021) analysed capacity bounds. Imani et al. (2019) applied HDC to classification tasks.

Knowledge distillation and model compression. Hinton et al. (2015) introduced knowledge distillation via soft targets. Romero et al. (2015) proposed FitNets for thin-deep networks. Jang et al. (2019) studied cross-architecture distillation. Our work transfers at the embedding level, not the logit level.

Neuro-symbolic integration. Garcez et al. (2019) surveyed neural-symbolic computing. Marra et al. (2020) combined logical constraints with neural networks. Hamilton et al. (2018) studied embedding-based knowledge graph completion. Our transplantation enables such embedding-based methods to interface with symbolic VSA operations.

8. Conclusion

We presented Model Transplantation for the NSCK architecture: a pipeline that projects dense pretrained embeddings into 10,240-bit binary hypervectors using random Gaussian projection followed by sign binarisation. Experiments confirmed Spearman $\rho > 0.995$, Recall@10 = 0.910–0.929, and cluster separation of 0.44 — sufficient for semantic nearest-neighbour query and VSA compositional operations. Round-trip fidelity is perfect (1.000) and projection is deterministic. Future work includes testing on real BERT/GPT embeddings, evaluating learned projections, and studying the interaction between transplanted codebooks and STDP-based SNN concept learning.

Acknowledgements

The author used AI coding assistants as iterative development and pair-programming tools during implementation of the NSCK codebase. All architectural decisions, experimental design, theoretical framing, and written content are the author’s own. This work was conducted independently, without institutional funding.

References

- Achlioptas, D. (2003). Database-friendly random projections: Johnson-Lindenstrauss with binary coins. *Journal of Computer and System Sciences*, 66(4), 671–687.
- Charikar, M. S. (2002). Similarity estimation techniques from rounding algorithms. *Proceedings of STOC*, 380–388.
- Dasgupta, S., & Gupta, A. (1999). *An elementary proof of the Johnson-Lindenstrauss lemma*. Technical Report TR-99-006, ICSI Berkeley.
- Frady, E. P., Kleyko, D., & Sommer, F. T. (2021). Variable binding for sparse distributed representations: Theory and applications. *IEEE Transactions on Neural Networks and Learning Systems*, 32(5), 2118–2132.
- Garcez, A. D., Gori, M., Lamb, L. C., Serafini, L., Spranger, M., & Tran, S. N. (2019). Neural-symbolic computing: An effective methodology for principled integration of machine learning and reasoning. *Journal of Applied Logics*, 6(4), 611–632.
- Gayler, R. W. (2004). Vector symbolic architectures answer Jackendoff’s challenges for cognitive neuroscience. *Proceedings of the ICCS/ASCS Joint International Conference on Cognitive Science*, 133–138.
- Gionis, A., Indyk, P., & Motwani, R. (1999). Similarity search in high dimensions via hashing. *Proceedings of VLDB*, 518–529.
- Hamilton, W., Leskovec, J., & Jurafsky, D. (2018). Diachronic embedding for facts in knowledge bases. *arXiv preprint arXiv:1812.10005*.
- Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*.
- Imani, M., Huang, C., Kong, D., & Rosing, T. (2019). Hierarchical hyperdimensional computing for energy efficient classification. *Proceedings of DAC*, 1–6.
- Indyk, P., & Motwani, R. (1998). Approximate nearest neighbors: Towards removing the curse of dimensionality. *Proceedings of STOC*, 604–613.
- Jang, Y., Lee, H., Hwang, S. J., & Shin, J. (2019). Learning what and where to transfer. *Proceedings of ICML*.
- Johnson, W. B., & Lindenstrauss, J. (1984). Extensions of Lipschitz mappings into a Hilbert space. *Contemporary Mathematics*, 26, 189–206.
- Kanerva, P. (2009). Hyperdimensional computing: An introduction to computing in distributed representation with high-dimensional random vectors. *Cognitive Computation*, 1(2), 139–159.

Li, P., Hastie, T. J., & Church, K. W. (2006). Very sparse random projections. *Proceedings of KDD*, 287–296.

Marra, G., Giannini, F., Diligenti, M., & Gori, M. (2020). Integrating learning and reasoning with deep logic models. *Proceedings of ECML-PKDD*, 517–532.

Plate, T. A. (2003). *Holographic Reduced Representations: Distributed representation for cognitive structures*. CSLI Publications, Stanford, CA.

Rachkovskij, D. A., & Kussul, E. M. (2001). Binding and normalization of binary sparse distributed representations by context-dependent thinning. *Neural Computation*, 13(2), 411–452.

Romero, A., Ballas, N., Kahou, S. E., Chassang, A., Gatta, C., & Bengio, Y. (2015). FitNets: Hints for thin deep nets. *Proceedings of ICLR*.

Vyas, R., & Bhattacharyya, C. (2022). Knowledge transfer via dense cross-layer mutual distillation. *arXiv preprint arXiv:2206.02840*.

Source code: `python/core/transplant/projector.py`, `python/core/transplant/validator.py`,
`python/core/vsa/vsa_embedding_bridge.py`. Experiments: `research/experiments/paper5_transplant_ben`
Results: `research/results/paper5_results.json`.